Zewail City
of Science, Technology and Innovation

CSAI School
CSAI 498 / CSAI 499

Pulse: Social Media and Marketing Multi-Agent Platform

"Intelligent Automation. Authentic Engagement. Real Results."

Team Members:

Student Name	ID	Program
Abdelrahman Elattar	202201353	DSAI
Abdelrahman Omar	202202254	DSAI
Rana Mahmoud	202201271	DSAI
Hager Saad	202201494	DSAI

Supervisor: Dr. Mohamed Maher

Semester: Fall 2025 / Spring 2026

Date of Submission: October 25, 2025

GitHub Repository

Contents

Abstract

Social media platforms have emerged as critical channels for businesses to establish brand presence, foster customer engagement, and derive consumer behavioral insights. However, managing multi-platform social media ecosystems demands substantial resources and frequently results in operational inconsistencies that undermine brand integrity. This project proposes **Pulse**, an autonomous AI-powered multi-agent platform designed to comprehensively automate social media management workflows. The system employs a sophisticated multi-agent architecture comprising four specialized agents: an AI Marketing Strategist for strategic oversight, a Social Media and Performance Analyst for content automation and analytics, an AI Customer Relationship Agent for intelligent engagement, and a Content and Market Authority Engine for autonomous content generation. Leveraging advanced natural language processing, large language models, and predictive analytics, Pulse orchestrates content scheduling, maintains brand-consistent audience interactions, and generates actionable performance insights through real-time data integration. The anticipated outcomes include significant operational efficiency gains, enhanced audience engagement metrics, improved response consistency across platforms, and substantial reduction in manual intervention requirements. This research demonstrates the practical application of multi-agent AI systems in addressing contemporary digital marketing challenges while maintaining scalability and adaptability.

1 Problem Statement & Motivation

1.1 The Problem

Managing social media is usually difficult for businesses, as it requires constant content creation, timely posting, audience engagement, and performance tracking. When handled manually, this process can lead to delays, missed opportunities, and inconsistent communication, ultimately damaging a brand's reputation and customer trust. Smaller businesses are particularly affected due to limited resources. We aim to address the challenge by developing an Agent workflow that automates scheduling, engagement, and performance analysis, helping marketing teams work more efficiently, maintaining a consistent brand presence, and ensuring customers receive timely, personalized interactions.

1.2 Motivation

The motivation for this project is to help in solving the challenges of managing social media by automating scheduling, managing audience engagement, and tracking performance metrics. **Pulse** represents our vision of keeping businesses connected to their audience's heartbeat—monitoring, responding, and adapting in real-time. This will benefit:

- **Marketing teams:** by reducing repetitive tasks and allowing them to focus on strategy and creativity.
- **Brands:** by maintaining a consistent, authentic voice across all platforms.
- **Consumers:** by providing faster, more relevant, and personalized interactions.

2 Proposed Solution

Pulse will function as an AI-powered assistant for social media and digital marketing, built around four core epics. Each epic represents a specialized agent designed to handle distinct but interconnected aspects of the marketing lifecycle. The platform's name reflects its core

capability: maintaining a constant pulse on brand presence, audience sentiment, and market dynamics.

2.1 Epic 1: AI Marketing Strategist

Goal:

The AI Marketing Strategist acts as the system's central intelligence, tackling the problem of uncoordinated oversight in marketing activities. Its primary goal is to monitor performance while also defining high-level objectives, ensuring brand consistency, and scattered execution across agents. For marketing directors, it provides data-driven strategies to optimize campaigns from the outset. For brand managers, it safeguards a consistent brand identity across all content. For business owners, it delivers continuous Key Performance Indicator monitoring and real-time recommendations, helping the organization adapt quickly to market changes and remain competitive.

Key Tasks:

1. Develop a knowledge base containing brand guidelines, historical campaign data, and audience personas.
2. Incorporate performance data by integrating with analytics and customer relationship management solutions.
3. Implement a decision-making engine that aligns campaign proposals with brand guidelines and key performance indicators.
4. Enable periodic, automated strategy reviews and optimization based on performance.
5. Build inter-agent communication protocols for delegating tasks to the Social Media, Content, and Ads agents.

2.2 Epic 2: Social Media & Performance Analyst Agent

Goal:

The system unifies social media automation and performance analytics into one intelligent agent that manages content and adapts strategy through real-time learning. It automates the full lifecycle planning, scheduling, engagement, and reporting while aggregating data from multiple sources to identify trends, forecast outcomes, and deliver actionable insights. The agent generates and schedules optimized content calendars, evaluates post-performance in real time, and refines future recommendations. It provides dashboards, predictive forecasts, and alerts on deviations, reducing manual effort and enabling faster, more strategic decision-making.

Key Tasks:

1. **Content Scheduling & Automation:** Build a scheduler that adapts posting times based on engagement data and trending topics.
2. **Brand Voice Consistency:** Fine-tune a conversational model using historical posts and audience interactions to ensure responses remain authentic and aligned with the brand.
3. **Performance Data Integration:** Engineer pipelines to ingest and normalize metrics from multiple platforms.
4. **Analytics & Visualization:** Create an integrated dashboard to display key performance indicators such as reach, impressions, CTR, conversions, and audience growth.
5. **Predictive Insights & Alerts:** Implement anomaly detection and forecasting to flag underperforming campaigns and highlight emerging trends.

6. **Automated Reporting:** Develop a reporting engine that generates daily, weekly, and monthly summaries with data-driven recommendations.

2.3 Epic 3: AI Customer Relationship Agent

Goal:

To deliver fast, consistent, and empathetic support by unifying all customer messages, voice notes, and communication channels under one AI-powered agent.

The AI Customer Support Agent is designed to resolve inquiries end-to-end across email, chat, social, voice, and in-app threads without human hand-off for the majority of cases. It will automatically classify incoming tickets, extract intent and sentiment from text and voice, and generate contextual replies that mirror brand tone. By pulling real-time data from the CRM, order system, and knowledge base, it provides personalized solutions, proactive updates, and relevant upsell answers. The agent continuously learns from every interaction, escalates only when necessary, and feeds conversation insights back to product and success teams to reduce future contact volume and improve the customer experience.

Key Tasks:

1. Securely integrate with CRM, telephony, social, and ticketing platforms.
2. Develop NLP models for multilingual intent, entity, and sentiment extraction from text and voice.
3. Build a reply-generation engine that respects brand voice, policy, and customer history.
4. Implement voice-to-text and text-to-voice pipelines for seamless voice-note handling.
5. Establish real-time knowledge-base retrieval and order-system actions for instant resolutions.
6. Create smart escalation and routing rules based on complexity, sentiment, and customer tier.
7. Deploy continuous-learning feedback loop with the AI Performance Analyst to track CSAT, FCR, deflection rate, and churn impact.

2.4 Epic 4: The Autonomous Content & Market Authority Engine

Business Goal (The "Why"):

To autonomously build and defend a dominant market position by establishing unassailable topical authority. This engine will capture high-intent organic traffic, systematically dismantle competitor content moats, and convert audience engagement into a measurable, revenue-generating pipeline.

Architectural Vision (The "How"):

A self-improving, multi-agent cognitive framework governed by a central Orchestrator. Specialized agents (e.g., Market Researcher, Content Strategist, Semantic Modeler, Generative Copywriter, SEO Optimizer, Performance Analyst) collaborate to autonomously model our domain, deconstruct user intent, execute complex, multi-quarter content campaigns, and refine their own strategies based on closed-loop, full-funnel performance data.

Core User Stories:

- **As a Chief Marketing Officer (CMO),** I want to set a single high-level goal like "Dominate the 'AI for E-commerce' market," so the Engine autonomously plans, executes, and reports on the entire multi-quarter content strategy to achieve this objective.
- **As the AI Marketing Strategist Agent,** I want to query this Engine for "a complete content plan to capture the 'high-value, at-risk' segment identified by the Customer Intelligence Agent," so it can generate a targeted retention campaign.

- **As a Head of Content,** I want the Engine to analyze the top 5 ranking articles for a high-intent keyword and generate a draft that is semantically richer and factually superior, so my team can shift from "first-draft writing" to "final-pass validation."
- **As the Content Orchestrator Agent,** upon publishing a new pillar page, I want to autonomously task the Social Media Agent to "atomize this asset into a 5-post social media campaign" and the Email Agent to "draft a newsletter segment," ensuring full asset syndication.
- **As the Performance Analyst Agent,** I want to detect when an article's conversion rate drops 20% post-SERP update, so I can autonomously task the SEO Optimizer Agent to re-analyze user intent and generate a content refresh brief.

Core Capabilities (The "What"):

- **Autonomous Opportunity Analysis & Strategic Planning:** This system proactively monitors the competitive landscape, SERP volatility, and internal data (from other agents like Customer Intelligence) to identify and prioritize high-ROI "content gaps." It autonomously generates complete, long-term content strategies (pillar pages, topic clusters) designed to achieve specific, high-level business goals (e.g., "Own the 'AI-driven finance' conversation").
- **Deep Semantic Modeling & Intent-Driven Generation:** Moving far beyond keywords, this system builds and maintains a deep semantic model of our domain. It autonomously deconstructs top-ranking competitor content to understand user intent, structural patterns, and argumentation flows. It then directs the Generative Agent to create semantically, factually, and structurally superior content that is precisely engineered to satisfy that intent.
- **Collaborative Content Syndication & Asset Atomization:** The system does not just "publish and pray." Upon finalizing a core asset (e.g., a pillar page), the Orchestrator autonomously tasks other agents (e.g., Social Media, Email) with a comprehensive syndication plan, atomizing the long-form content into a cascade of derivative assets (social posts, newsletter snippets, ad copy), maximizing the ROI of each strategic effort.
- **Full-Funnel Performance Optimization & Learning Loop:** This engine autonomously tracks content performance from initial ranking and traffic (via Search Console) to deep engagement and conversion (via Analytics/CRM). This data feeds a reinforcement learning loop that continuously refines generation models, keyword prioritization, and its understanding of what content converts, ensuring the system becomes more valuable and effective with every cycle.

Key Performance Indicators (The "Value"):

- Increase in Share of Voice (SOV) for target, high-intent keyword clusters.
- Growth in Content-Attributable Revenue Pipeline (MQLs, SQLs, and closed-won opportunities).
- Reduction in Content-to-Pipeline Velocity (time from strategic gap identification to a revenue-generating asset).
- Improvement in Cost-per-Acquisition (CPA) from all organic channels.

3 Project Scope

3.1 In Scope

- AI-driven content generation and scheduling across Facebook/Instagram, LinkedIn, and X.
- NLP-powered response system for comments and direct messages.
- Integration with official APIs for posting and retrieving analytics data.
- Performance dashboards with detailed metrics and actionable insights.

3.2 Out of Scope

- Paid advertising campaign management.
- Non-social digital marketing channels (e.g., search ads, offline marketing).
- Multi-Language Support: All functionality is limited to the English language.

4 High-Level Timeline

Phase	Description	Duration (weeks)	Deliverables
Fall 2025 Semester			
Research & Requirement Analysis	Project proposal development, literature review, requirement gathering, and feasibility analysis	4 weeks	Project Proposal (Week 4)
Research & Requirement Analysis (cont.)	Learning phase: research and skill building in LLMs, NLP, multi-agent systems, and API integration	6 weeks	Research Summary, Technology Evaluation Report (Week 10, Midterm Week 8)
Design & Planning	System architecture design, database schema, agent interaction protocols, and API specifications	3 weeks	Initial System Design Document (Week 11-13)
Design & Planning (cont.)	Finalize design documentation and prepare comprehensive Part-1 report	1 week	Final Part-1 Report (Week 14)
Spring 2026 Semester			
Implementation Part 1	Backend development, database setup, core API integrations, and initial agent logic implementation	4 weeks	Implementation Progress Report 1 (Week 4)
Implementation Part 1 (cont.)	Frontend dashboard development, basic automation features, and MVP deployment	2 weeks	MVP Delivery (Week 6)
Implementation Part 1 (cont.)	Advanced features development, agent orchestration, and analytics engine implementation	2 weeks	Implementation Progress Report 2 (Week 8)
Testing & Evaluation	Comprehensive system testing, performance optimization, bug fixes, and user acceptance testing	5 weeks	Final Report + User Guide (Week 13)
Testing & Evaluation (cont.)	Final presentation preparation, system demonstration, and repository finalization	1 week	Demo + Oral Defense, GitHub Repository Final Submission (Week 14)

Table 1: Project Timeline with Deliverables

5 Technology Stack & Theoretical Basis

5.1 Programming Languages & Frameworks

Backend:

- **Python (FastAPI):** Selected for its asynchronous capabilities, which are essential for real-time message processing and handling concurrent requests efficiently.

Frontend:

- **React:** Chosen for building a responsive web dashboard for administrators, providing modern UI components and excellent user experience.

5.2 Database

- **MongoDB:** Selected for its flexible schema, making it ideal for storing diverse message data, moderation logs, and analytics information with varying structures.

5.3 AI Engine

- **Pre-trained Large Language Model:** Core AI component for contextual understanding and response generation.
- **CrewAI:** Framework for orchestrating multi-agent systems and coordinating specialized agents.
- **AutoGen:** Multi-agent conversation framework for building sophisticated agent interactions and workflows.

5.4 APIs & Integrations

- **GA4 Data API:** Google Analytics 4 integration for comprehensive web analytics and performance tracking.
- **Reddit API:** For social listening, trend monitoring, and community engagement analysis.
- **GDELT:** Global Database of Events, Language, and Tone for real-time global news and event monitoring.
- **NewsAPI:** Real-time news aggregation for content inspiration and trend detection.

5.5 DSAI Program-Specific Focus

As Data Science and Artificial Intelligence students, our technical approach emphasizes:

Data Lifecycle Management:

- **Data Collection:** Designing robust ETL pipelines for ingesting multi-platform social media data with automated validation and quality control.
- **Data Storage:** Implementing efficient NoSQL database schemas optimized for social media analytics and historical data warehousing.
- **Data Processing:** Developing scalable data processing workflows for batch analytics and real-time stream processing.
- **Data Analysis:** Applying statistical methods, time-series analysis, and machine learning techniques to extract actionable business insights.

AI/ML Techniques:

- Fine-tuning large language models for domain-specific brand voice consistency.
- Implementing NLP pipelines for sentiment analysis, intent classification, and entity extraction.
- Developing predictive analytics models for engagement forecasting and anomaly detection.
- Creating multi-agent coordination systems using CrewAI and AutoGen frameworks.

Stakeholder-Centric Analytics:

- Designing intuitive dashboards that translate complex analytics into actionable business insights.
- Implementing real-time performance monitoring with automated alerting systems.
- Creating comprehensive reporting engines for strategic decision support.

6 Success Metrics & Evaluation Plan

6.1 Performance Metrics

System Performance:

- **Response Time:** Average API response time under 200ms for real-time operations
- **Content Generation Speed:** Automated post creation completed within 5 seconds
- **System Uptime:** Maintain 99% availability during operational hours
- **Processing Throughput:** Handle minimum 1,000 social media interactions per hour

AI Model Accuracy:

- **Intent Classification:** Achieve greater than 85% accuracy on test datasets
- **Sentiment Analysis:** Achieve greater than 80% F1-score across sentiment classes
- **Brand Voice Consistency:** Achieve greater than 80% human evaluator approval rating
- **Response Relevance:** Generate greater than 85% contextually appropriate responses

6.2 Business Impact Metrics

Efficiency Gains:

- **Time Savings:** Reduce manual content scheduling time by at least 60%
- **Response Time Improvement:** Decrease average customer inquiry response time by 50%
- **Content Volume Increase:** Enable 40% increase in published content without additional resources

Engagement Metrics:

- **Engagement Rate:** Improve average post engagement rate by 20%
- **Response Coverage:** Achieve greater than 90% response rate to comments and messages
- **Audience Growth:** Support consistent monthly follower growth through optimized posting

6.3 Quality Metrics

Usability:

- **User Satisfaction:** Achieve System Usability Scale (SUS) score above 70
- **Ease of Use:** Enable new users to become proficient in core features within 30 minutes
- **Dashboard Clarity:** Achieve greater than 80% user comprehension of analytics visualizations

Reliability & Security:

- **Error Rate:** Maintain less than 3% failed API calls or system operations
- **Data Security:** Implement encryption and secure authentication for all integrations
- **Audit Trail:** Maintain complete logging of all automated actions for accountability

6.4 Evaluation Methodology

Testing Approach:

1. **Unit Testing:** Comprehensive automated tests for individual components with high code coverage
2. **Integration Testing:** Validate seamless inter-agent communication and external API integrations
3. **Performance Testing:** Conduct load testing with simulated high-volume scenarios
4. **User Acceptance Testing:** Deploy pilot system with selected test users for real-world validation
5. **A/B Testing:** Compare automated content performance against baseline manual processes

Data Collection Methods:

- Application performance logs and monitoring dashboards
- User interaction analytics and behavior tracking
- Post-deployment user surveys and feedback collection
- Social media platform analytics via official APIs
- Comparative analysis of pre- and post-deployment KPIs

7 Team Roles & Responsibilities

Team Member	Program	Primary Role	Technical Contribution (DSAI Focus)
Abdelrahman Elattar	DSAI	AI/ML Lead & Project Coordinator	Lead LLM fine-tuning and agent orchestration using CrewAI/AutoGen. Develop NLP pipelines for intent classification, sentiment analysis, and brand voice consistency models. Coordinate project milestones and advisor communication.
Abdelrahman Omar	DSAI	Data Engineering & Backend Lead	Design and implement data collection pipelines, ETL processes, and database architecture. Build backend APIs using FastAPI and manage real-time data streaming infrastructure for multi-platform integration.
Rana Mahmoud	DSAI	Analytics & Visualization Lead	Develop predictive analytics models, anomaly detection algorithms, and performance forecasting systems. Create interactive React dashboards with comprehensive data visualizations for actionable insights.
Hager Saad	DSAI	Integration & Deployment Lead	Implement API integrations with social media platforms (Facebook, Instagram, LinkedIn, X), GA4, Reddit, GDELT, and NewsAPI. Manage deployment infrastructure, testing frameworks, and system optimization.

Table 2: Team Roles and DSAI-Specific Contributions

7.1 Shared Responsibilities

All team members collectively contribute to:

- Code review and quality assurance
- Comprehensive documentation (technical, user guides, API documentation)
- Testing activities (unit, integration, user acceptance)
- Research, literature review, and technology evaluation
- Presentation preparation and deliverable submissions

7.2 Communication & Collaboration Plan

Internal Team Communication:

- **Platform:** Slack workspace for daily communication, file sharing, and quick coordination
- **Channels:** Organized channels including #general, #technical-discussion, #resources, #meetings

- **Response Commitment:** Team members respond to urgent messages within 4 hours during weekdays

Team Meetings:

- **Frequency:** Twice weekly (Sundays and Wednesdays)
- **Platform:** Google Meet for virtual collaboration
- **Duration:** 60-90 minutes per session
- **Agenda Structure:**
  - Progress updates and blockers (15 minutes)
  - Technical discussions and problem-solving (30 minutes)
  - Sprint planning and task assignments (20 minutes)
  - Action items and next steps (10 minutes)
- **Documentation:** Meeting minutes recorded with action items tracked in shared documents

Advisor Meetings:

- **Frequency:** Bi-weekly meetings with Dr. Mohamed Maher
- **Platform:** Google Meet or in-person as scheduled
- **Preparation:** Progress reports submitted 24 hours before meetings

Documentation & Development Platforms:

- **GitHub:** Version control, code repository, issue tracking, and project management Repository: https://github.com/el3ttar3/pulse-multi-agent-platform
- **Google Drive:** Shared documentation, reports, and presentations
- **Notion/Confluence:** Project wiki, resource library, and knowledge base for Pulse platform

Development Methodology:

- **Approach:** Agile methodology with 2-week sprint cycles
- **Task Management:** GitHub Projects for sprint planning, backlog management, and progress tracking
- **Code Standards:** Established coding conventions, documentation requirements, and review processes

8 References

1. T. Brown et al., "Language Models are Few-Shot Learners," *Advances in Neural Information Processing Systems*, vol. 33, pp. 1877-1901, 2020.
2. J. Wei et al., "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models," *Advances in Neural Information Processing Systems*, vol. 35, 2022.
3. A. Vaswani et al., "Attention Is All You Need," *Advances in Neural Information Processing Systems*, vol. 30, 2017.
4. J. Devlin et al., "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding," *Proceedings of NAACL-HLT*, pp. 4171-4186, 2019.
5. CrewAI Documentation, "Multi-Agent Orchestration Framework," [Online]. Available: https://docs.crewai.
6. Microsoft AutoGen Documentation, "Building Multi-Agent Applications," [Online]. Available: https://microsoft.github.io/autogen/
7. Facebook for Developers, "Graph API Documentation," Meta Platforms Inc., [Online]. Available: https://developers.facebook.com/docs/graph-api
8. LinkedIn Developers, "LinkedIn API Documentation," [Online]. Available: https://docs.microsoft.com/en-us/linkedin/
9. X Developer Platform, "Twitter API v2 Documentation," [Online]. Available: https://developer.twitter.com
10. Google Analytics, "GA4 Data API Documentation," Google LLC, [Online]. Available: https://developers.google.com/analytics/devguides/reporting/data/v1
11. Reddit Inc., "Reddit API Documentation," [Online]. Available: https://www.reddit.com/dev/api/
12. NewsAPI, "News API Documentation," [Online]. Available: https://newsapi.org/docs
13. GDELT Project, "The GDELT Project Documentation," [Online]. Available: https://www.gdeltproject.org/
14. S. Ramirez, "FastAPI: Modern Web Framework for Building APIs," [Online]. Available: https://fastapi.tiangolo.com/
15. MongoDB Inc., "MongoDB Documentation," [Online]. Available: https://www.mongodb.com/docs/
16. Meta AI, "React Documentation," [Online]. Available: https://react.dev/
17. Y. Liu et al., "RoBERTa: A Robustly Optimized BERT Pretraining Approach," *arXiv preprint arXiv:1907.11692*, 2019.
18. H. Touvron et al., "LLaMA: Open and Efficient Foundation Language Models," *arXiv preprint arXiv:2302.13971*, 2023.
19. S. Bird, E. Klein, and E. Loper, *Natural Language Processing with Python*, O'Reilly Media Inc., 2009.
20. M. Honnibal and I. Montani, "spaCy 2: Natural Language Understanding with Bloom Embeddings, Convolutional Neural Networks and Incremental Parsing," 2017.
21. F. Pedregosa et al., "Scikit-learn: Machine Learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825-2830, 2011.

22. A. Radford et al., "Improving Language Understanding by Generative Pre-Training," OpenAI, 2018.

23. OpenAI, "GPT-4 Technical Report," *arXiv preprint arXiv:2303.08774*, 2023.

24. S. Bubeck et al., "Sparks of Artificial General Intelligence: Early Experiments with GPT-4," *arXiv preprint arXiv:2303.12712*, 2023.

25. W. X. Zhao et al., "A Survey of Large Language Models," *arXiv preprint arXiv:2303.18223*, 2023.